

---

# COMP9311: DATABASE SYSTEMS


Term 1 2024

Week 3 – SQL

By Xiaoyang Wang, CSE UNSW

---

*Disclaimer: the course materials are sourced from previous offerings of  
COMP9311 and COMP3311*

A solid orange horizontal bar at the bottom of the slide.

# Summary Week 2

---

## **Relational Model**

- Relations, Tuples, Attributes
- Integrity Constraints
- ER to Relational Mapping

## **Relational Algebra**

- Select, Project, Union, Intersection, Difference, Cartesian Product, Join, Divide, Rename
- Basic Operators vs Extended Operators
- Aggregation

# SQL-99

---

- SQL = **Structured Query Language** (pronounced “sequel”).
- Developed at IBM (San Jose Lab) during the 1970's, and standardised during the 1980's.
- A **standard** language for querying and manipulating relational DBMSs.
- Interactive via GUI or command line or embedded in programs.
- **Declarative**, based on relational algebra.

# SQL in Relational DBMS

---

In relational databases, what does SQL do?

## **A data definition language (DDL)**

- CREATE TABLE, DROP TABLE, ...

## **A data manipulation language (DML)**

- SELECT: keywords relating to select, e.g., GROUP BY, HAVING, ORDER BY...
- INSERT, DELETE, UPDATE, ALTER, ...

## **Other Commands**

- indexes, constraints, views, triggers, transactions, authorization, ...

# Sample Database

---

To illustrate the features of SQL, we use a small example database below:

Beers(*name*, manf)

Bars(*name*, addr, license)

Drinkers(*name*, addr, phone)

Likes(*drinker*, *beer*)

Sells(*bar*, *beer*, price)

Frequents(*drinker*, *bar*)

keys are in ***italic*** font and highlighted by **underscore**.

# Sample Database (cont)

---

## Bars:

Name	Addr	License
Australia Hotel	The Rocks	123456
Coogee Bay Hotel	Coogee	966500
Lord Nelson	The Rocks	123888
Marble Bar	Sydney	122123
Regent Hotel	Kingsford	987654
Royal Hotel	Randwick	938500

## Drinkers:

Name	Addr	Phone
Adam	Randwick	9385-4444
Gernot	Newtown	9415-3378
John	Clovelly	9665-1234
Justin	Mosman	9845-4321

# Sample Database (cont)

---

**Beers:**

Name	Manf
80/-	Caledonian
Bigfoot Barley Wine	Sierra Nevada
Burraborang Bock	George IV Inn
Crown Lager	Carlton
Fosters Lager	Carlton
Invalid Stout	Carlton
Melbourne Bitter	Carlton
New	Toohey's
Old	Toohey's
Old Admiral	Lord Nelson
Pale Ale	Sierra Nevada
Premium Lager	Cascade
Red	Toohey's
Sheaf Stout	Toohey's
Sparkling Ale	Cooper's
Stout	Cooper's
Three Sheets	Lord Nelson
Victoria Bitter	Carlton

# Sample Database (cont)

## Frequents:

Drinker	Bar
Adam	Coogee Bay Hotel
Gernot	Lord Nelson
John	Coogee Bay Hotel
John	Lord Nelson
John	Australia Hotel
Justin	Regent Hotel
Justin	Marble Bar

## Likes:

Drinker	Beer
Adam	Crown Lager
Adam	Fosters Lager
Adam	New
Gernot	Premium Lager
Gernot	Sparkling Ale
John	80/-
John	Bigfoot Barley Wine
John	Pale Ale
John	Three Sheets
Justin	Sparkling Ale
Justin	Victoria Bitter



# Sample Database (cont)

**Sells:**

Bar	Beer	Price
Australia Hotel	Burraborang Bock	3.5
Coogee Bay Hotel	New	2.25
Coogee Bay Hotel	Old	2.5
Coogee Bay Hotel	Sparkling Ale	2.8
Coogee Bay Hotel	Victoria Bitter	2.3
Lord Nelson	Three Sheets	3.75
Lord Nelson	Old Admiral	3.75
Marble Bar	New	2.8
Marble Bar	Old	2.8
Marble Bar	Victoria Bitter	2.8
Regent Hotel	New	2.2
Regent Hotel	Victoria Bitter	2.2
Royal Hotel	New	2.3
Royal Hotel	Old	2.3
Royal Hotel	Victoria Bitter	2.3

Example:

**Beers:**

Name	Manf
80/-	Caledonian
Bigfoot Barley Wine	Sierra Nevada
Burraborang Bock	George IV Inn
Crown Lager	Carlton
Fosters Lager	Carlton
Invalid Stout	Carlton
Melbourne Bitter	Carlton
New	Toohey's
Old	Toohey's
Old Admiral	Lord Nelson
Pale Ale	Sierra Nevada
Premium Lager	Cascade
Red	Toohey's
Sheaf Stout	Toohey's
Sparkling Ale	Cooper's
Stout	Cooper's
Three Sheets	Lord Nelson
Victoria Bitter	Carlton

SQL Queries: What beers are made by Toohey's?"

Example:

**Beers:**

Name	Manf
80/-	Caledonian
Bigfoot Barley Wine	Sierra Nevada
Burraborang Bock	George IV Inn
Crown Lager	Carlton
Fosters Lager	Carlton
Invalid Stout	Carlton
Melbourne Bitter	Carlton
New	Toohey's
Old	Toohey's
Old Admiral	Lord Nelson
Pale Ale	Sierra Nevada
Premium Lager	Cascade
Red	Toohey's
Sheaf Stout	Toohey's
Sparkling Ale	Cooper's
Stout	Cooper's
Three Sheets	Lord Nelson
Victoria Bitter	Carlton

SQL Queries: What beers are made by Toohey's?"

**SELECT Name FROM Beers WHERE Manf = 'Toohey's';**

# SQL Queries

---

To answer the question “What beers are made by Toohey’s?”, we could ask:

**SELECT Name FROM Beers WHERE Manf = ‘Toohey’s’;**

This gives a subset of the Beers relation, displayed as:

Name

-----

New

Old

Red

Sheaf Stout

Quotes are escaped by doubling them (‘ ‘)

# Basic SELECT Structure

---

To retrieve information from a database, there is a basic query structure, known as the **select** statement

SELECT <Attribute list>

FROM <Table list>

WHERE <Condition>

- <attribute list>: list of attributes
- <table list>: list of relations
- <condition>: list of conditions (Boolean expression)

SELECT statement is also known as a **select-from-where block**. The result of this statement is a table, which is typically displayed on output.

The SELECT statement contains the functionality of **select**, **project** and **join** from the relational algebra.

# SQL Identifiers

---

Names are used to identify objects such as tables, attributes, views, ...

Identifiers in SQL use similar conventions to common programming languages:

- a sequence of **alpha-numeric**s, starting with an alphabetic
- **not** case-sensitive
- reserve word **disallowed**, ...
- Whether you can use sign (@), dollar sign (\$), number sign (#), or underscore (\_) depends on the standard

# SQL Keywords

---

Some of the frequently-used ones

- ALTER AND CREATE
- FROM INSERT NOT OR
- SELECT TABLE WHERE

For PostgreSQL Keywords see

<https://www.postgresql.org/docs/current/sql-keywords-appendix.html> .

# SQL Data Types

---

All attributes in SQL relations have **domain** specified.

SQL supports a small set of useful built-in data types: strings, numbers, dates, bit-strings.

Self defined data type is allowed in PostgreSQL.

Various type conversions are available

- date to string, string to date, integer to real ...
- applied automatically “where they make sense”



# SQL Data Types (cont.)

---

- Basic domain (type) checking is performed automatically
- Constraints can be used to “enforce” more complex domain membership conditions
- The NULL value is a member of all data types

# SQL Data Types (cont.)

---

Comparison operators are defined on all types

<      >      <=      >=      =      !=

Boolean operators **AND**, **OR**, **NOT** are available within WHERE expressions to combine results of comparisons

Comparison against NULL yields FALSE

Can explicitly test for NULL using:

***attr* IS NULL**

***attr* IS NOT NULL**

Most data types also have type-specific operations available (e.g., arithmetic for numbers)

Which operations are actually applied depends on the implementation

# Data Types - Numeric

---

Some options for specifying the attributes for holding numeric values:

If you need integers

- **smallint** (2 byte integer)
- **int** (4 byte integer)
- **bigint** (8 byte integer)

If you need real numbers

- **real** (4 byte floating point)
- **double** (8 byte floating point)
- **numeric (<precision> , <scale>)**
  - <precision>: specify significant digits in the whole number
  - <scale>: specify digits after the decimal point

# Data Types – String Literal

---

Example of a string literal: 'John'

A string literal is **a sequence of zero or more characters**

In SQL, you specify a literal by enclosing it in single quotes.

Two kinds of string literals are available:

- **CHAR(n)** **n length**, left-justified blank-padded
- **VARCHAR(n)** **can be between 0 and n length**, no padding

String literals are case sensitive: 'John' != 'JOHN'

# String Operators

---

- `string || string` ... concatenate two strings
- `LENGTH (string)` ... return length of string
- `SUBSTR (string, start, length)` ... extract chars from within string

Example:

- `'Post' || 'greSQL' -> PostgreSQL`
- `substring('Thomas' from 2 for 3) -> hom`

# SQL Like Operator

---

**str LIKE pattern** ... matches string to pattern

Two kinds of string **pattern-matching**

- The symbol \_ (underscore) matches any single characters
- The symbol % (percent) matches zero or more characters

Practice

- |                       |                                |
|-----------------------|--------------------------------|
| ○ String LIKE 'Ja%'   | Strings beginning with 'Ja'    |
| ○ String LIKE '_i%'   | Strings with 'i' as 2nd letter |
| ○ String LIKE '%o%o%' | Strings contains two 'o's      |

# SQL Dates

---

Dates are simply specially-formatted strings, with a range of operations to implement date semantics.

Format is typically **DD-Mon-YYYY**, e.g., '18-Aug-1998'

Accepts other formats

Comparison operators implement before (<) and after (>).

**(start1, end1) OVERLAPS (start2, end2)**

- This expression yields true when two time periods (defined by their endpoints) overlap, false when they do not overlap.
- `SELECT (DATE '2001-02-16', DATE '2001-12-21') OVERLAPS (DATE '2001-10-30', DATE '2002-10-30');` -> Result: true

# Converting Data Types

---

Conversions between data types are an important skill to know.

E.g., division of one integer with an integer

Various type conversions are available:

- integer to real ...
- string to integer ...

SQL supports a small set of useful built-in data types, e.g., numbers, strings, dates...

- You can define your own type in SQL



# Tuple and Set Literals

---

Tuple and set constants are both written as:

**(val1, val2, val3, ... )**

The correct interpretation is worked out from the context.

Examples:

```
Student(stude#, name, course)
( 2177364, 'Jack Smith', 'BSc') -- tuple literal
```

```
SELECT name
FROM Employees
WHERE job IN ('Lecturer', 'Tutor', 'Professor'); -- set literal
```

# Querying a Single Relation

---

Formal semantics (relational algebra):

- start with relation  $R$  in FROM clause
- apply  $\sigma$  using Condition in WHERE clause
- apply  $\pi$  using Attributes in SELECT clause

SELECT Attributes

FROM  $R$

WHERE Conditions

# Querying a Single Relation

---

Operationally, we think in terms of a tuple variable ranging over all tuples of the relation.

## **Operational semantics of SQL SELECT (single relation)**

FOR EACH tuple **T** in R DO

    check whether T satisfies the condition in the WHERE clause

    IF it does THEN

        print the attributes of T that are  
        specified in the SELECT clause

    END

END

# Projection by SELECT Clause

---

The **select** clause lists the attributes desired in the result of a query

- corresponds to the **projection** operation of the relational algebra

Example: Give all the names of all drinkers

```
SELECT Name  
FROM Drinkers;
```

## Drinkers:

Name	Addr	Phone
Adam	Randwick	9385-4444
Gernot	Newtown	9415-3378
John	Clovelly	9665-1234
Justin	Mosman	9845-4321

Note: FROM is always necessary with SELECT, whereas WHERE is optional.

# Projection by SELECT Clause

---

The **select** clause lists the attributes desired in the result of a query

- corresponds to the **projection** operation of the relational algebra

Example: Give all the names of all drinkers

**SELECT Name**  
**FROM Drinkers;**

**Drinkers:**

Name	Addr	Phone
Adam	Randwick	9385-4444
Gernot	Newtown	9415-3378
John	Clovelly	9665-1234
Justin	Mosman	9845-4321

Note: FROM is always necessary with SELECT, whereas WHERE is optional.

# Projection by SELECT Clause

Example: Give me both names and addresses of drinkers

SELECT Name, Addr  
FROM Drinkers;

**Drinkers:**

Name	Addr	Phone
Adam	Randwick	9385-4444
Gernot	Newtown	9415-3378
John	Clovelly	9665-1234
Justin	Mosman	9845-4321

An asterisk in the select clause denotes “all attributes”

SELECT \*  
FROM Drinkers;

**Drinkers:**

Name	Addr	Phone
Adam	Randwick	9385-4444
Gernot	Newtown	9415-3378
John	Clovelly	9665-1234
Justin	Mosman	9845-4321

# DISTINCT

---

SQL allows duplicates in relations and in query results.

- allows a table to have two or more tuples that are identical in all their attribute values.

In general, an SQL table can be a simple set of tuples, or a multiset of tuples.

Set: {a, b, c}

Multiset: {a, a, b, b, c, a, a, b, c, c ...}

# DISTINCT

---

To eliminate duplicates in the query results, insert the keyword **distinct** after select.

Example: Find the names of all departments and remove duplicates.

**SELECT DISTINCT** dept\_name from instructor

<i>ID</i>	<i>name</i>	<i>dept_name</i>	<i>salary</i>
22222	Einstein	Physics	95000
12121	Wu	Finance	90000
32343	El Said	History	60000
45565	Katz	Comp. Sci.	75000
98345	Kim	Elec. Eng.	80000
76766	Crick	Biology	72000
10101	Srinivasan	Comp. Sci.	65000
58583	Califieri	History	62000
83821	Brandt	Comp. Sci.	92000
15151	Mozart	Music	40000
33456	Gold	Physics	87000
76543	Singh	Finance	80000

(a) The *instructor* table

Question: When duplicates are useful?



# Selection by Where Clause

Find the beers manufactured by Toohey's

```
SELECT Name  
FROM Beers  
WHERE Manf = 'Toohey's';
```

**Beers:**

Name	Manf
80/-	Caledonian
Premium Lager	Cascade
Red	Toohey's
Sheaf Stout	Toohey's
Sparkling Ale	Cooper's
Victoria Bitter	Carlton

The “typical” SELECT query:

```
SELECT a1, a2, a3  
FROM Rel  
WHERE Cond
```

This corresponds to select followed by project

$$\pi_{\{a1,a2,a3\}}(\sigma_{\text{Cond}}(\text{Rel}))$$

# Example

Find the price that Regent Hotel charges for  
New

```
SELECT price
FROM Sells
WHERE bar = 'Regent Hotel' AND beer = 'New';
```

```
PRICE
-----
2.2
```

The condition can be an arbitrarily complex  
boolean-valued expression using the  
operators mentioned previously.

Bar	Beer	Price
Australia Hotel	Burraborang Bock	3.5
Coogee Bay Hotel	New	2.25
Coogee Bay Hotel	Old	2.5
Coogee Bay Hotel	Sparkling Ale	2.8
Coogee Bay Hotel	Victoria Bitter	2.3
Lord Nelson	Three Sheets	3.75
Lord Nelson	Old Admiral	3.75
Marble Bar	New	2.8
Marble Bar	Old	2.8
Marble Bar	Victoria Bitter	2.8
Regent Hotel	New	2.2
Regent Hotel	Victoria Bitter	2.2
Royal Hotel	New	2.3
Royal Hotel	Old	2.3
Royal Hotel	Victoria Bitter	2.3

# Null in SQL

---

What happens when the condition makes a comparison with a null value?

**Comparisons** with null returns unknown

- Example:  $5 < \text{null}$ ,  $\text{null} <> \text{null}$ ,  $\text{null} = \text{null}$

Three-valued logic using the truth value unknown

- **OR**: (**unknown** or true) = true,  
(**unknown** or false) = **unknown**,  
(**unknown** or **unknown**) = **unknown**
- **AND**: (true and **unknown**) = **unknown**,  
(false and **unknown**) = false,  
(**unknown** and **unknown**) = **unknown**
- **NOT**: (not **unknown**) = **unknown**
- “P is unknown” evaluates to true if predicate P evaluates as unknown

Result of where clause predicate is treated as **false** if it evaluates as unknown

# Example: Null Values

```
SELECT A3
FROM R
WHERE A1 + 5 > A2 and A4 = 'x'
```

A1	A2	A3	A4
5	9	alpha	x
	4	beta	
2	4	gamma	
3		delta	x

When it evaluates the second tuple:

- **Null + 5 -> Null** (for A1 + 5)
- **Null > 4 -> unknown** (for A1 + 5 > A2)
- **Null = 'x' -> unknown** (for A4 = 'x')
- **unknown and unknown -> unknown** (for A1 + 5 > A2 and A4 = 'x')

What about the following?

```
select A3
from R
where (A1 + 5 > A2 and A4 = 'x') is unknown
```

# Renaming via AS

---

In relational algebra we have the renaming operator  $\rho$  to avoid name clashes.

Example:  $\rho_{\text{Beers}(\text{Brand}, \text{Brewer})}(\text{Beers})$

Gives a new relation, with same data as Beers, but with attribute names changed.

**SQL provides AS to achieve this and it is used in the SELECT part.**

# Renaming via as (cont.)

---

Example:

- Beers(name, manf)

**SELECT name AS Brand, manf AS Brewer FROM Beers;**

BRAND	BREWER
-----	-----
80/-	Caledonian
Bigfoot Barley Wine	Sierra Nevada
Burraborang Bock	George IV Inn
Crown Lager	Carlton
Fosters Lager	Carlton
Invalid Stout	Carlton
...	

# Expressions as Values in Columns

---

**AS** can also be used to introduce computed values

Example:

Sells(bar, beer, price)

```
SELECT bar, beer, price*120 AS PriceInYen  
FROM Sells;
```

BAR	BEER	PRICEINYEN
-----	-----	-----
Australia Hotel	Burraborang Bock	420
Coogee Bay Hotel	New	270
Coogee Bay Hotel	Old	300
Coogee Bay Hotel	Sparkling Ale	336
Coogee Bay Hotel	Victoria Bitter	276
...		

**Just Display but no change to the database**

# Inserting Text in Result Table

Trick: to put text in output columns, use constant expression with AS.

Example:

Likes(drinker, beer)

```
SELECT drinker, 'likes Cooper''s' AS WhoLikes
FROM Likes
WHERE beer = 'Sparkling Ale';
```

DRINKER	WHOLIKES
-----	-----
Gernot	likes Cooper's
Justin	likes Cooper's

Drinker	Beer
Adam	Crown Lager
Adam	Fosters Lager
Adam	New
Gernot	Premium Lager
Gernot	Sparkling Ale
John	80/-
John	Bigfoot Barley Wine
John	Pale Ale
John	Three Sheets
Justin	Sparkling Ale
Justin	Victoria Bitter



# Querying Multi-relations

Question: Find the brewers whose beers John likes?

**Likes:**

Drinker	Beer
Adam	Crown Lager
Adam	Fosters Lager
Adam	New
Gernot	Premium Lager
Gernot	Sparkling Ale
John	80/-
John	Bigfoot Barley Wine
John	Pale Ale
John	Three Sheets
Justin	Sparkling Ale
Justin	Victoria Bitter

**Beers:**

Name	Manf
80/-	Caledonian
Bigfoot Barley Wine	Sierra Nevada
Burraborang Bock	George IV Inn
Crown Lager	Carlton
Fosters Lager	Carlton
Invalid Stout	Carlton
Melbourne Bitter	Carlton
New	Toohey's
Old	Toohey's
Old Admiral	Lord Nelson
Pale Ale	Sierra Nevada
Premium Lager	Cascade
Red	Toohey's
Sheaf Stout	Toohey's
Sparkling Ale	Cooper's
Stout	Cooper's
Three Sheets	Lord Nelson
Victoria Bitter	Carlton

# Querying Multi-relations

Question: Find the brewers whose beers John likes?

**Likes:**

Drinker	Beer
Adam	Crown Lager
Adam	Fosters Lager
Adam	New
Gernot	Premium Lager
Gernot	Sparkling Ale
John	80/-
John	Bigfoot Barley Wine
John	Pale Ale
John	Three Sheets
Justin	Sparkling Ale
Justin	Victoria Bitter

**Beers:**

Name	Manf
80/-	Caledonian
Bigfoot Barley Wine	Sierra Nevada
Burraborang Bock	George IV Inn
Crown Lager	Carlton
Fosters Lager	Carlton
Invalid Stout	Carlton
Melbourne Bitter	Carlton
New	Toohey's
Old	Toohey's
Old Admiral	Lord Nelson
Pale Ale	Sierra Nevada
Premium Lager	Cascade
Red	Toohey's
Sheaf Stout	Toohey's
Sparkling Ale	Cooper's
Stout	Cooper's
Three Sheets	Lord Nelson
Victoria Bitter	Carlton

# Querying Multi-relations

---

Example: Find the brewers whose beers John likes

- Likes(drinker, beer)
- Beers(name, manf)

```
SELECT Manf
FROM Likes, Beers
WHERE drinker = 'John' AND beer = name
```

MANF

-----

Caledonian

Sierra Nevada

Sierra Nevada

Lord Nelson

Note: could eliminate the duplicates by using DISTINCT

Relation algebra:  $\pi_{\text{manf}}(\sigma_{\text{drinker}='john'}(\text{Likes} \bowtie_{\text{beer=name}} \text{Beers}))$

# The From Clause

---

For SQL SELECT statement on several relations:

**SELECT Attributes**

**FROM R1, R2, ...**

**WHERE Condition**

Formal semantics (relational algebra):

- start with **Cartesian Product**  $R1 \times R2 \times \dots$  in FROM clause
- apply  $\sigma$  using Condition in WHERE clause
- apply  $\pi$  using Attributes in SELECT clause

# Querying Multi-relations

---

## Operational semantics of SELECT (multi-relations):

```
FOR EACH tuple T1 in R1 DO
  FOR EACH tuple T2 in R2 DO
    ...
    check WHERE condition for current
    assignment of T1, T2, ... vars
    IF holds THEN
      print attributes of T1, T2, ...
      specified in SELECT
    END
  ...
END
END
```

**For efficiency reasons, it is not implemented in this way!**

# Cartesian Product

instructor

ID	name	dept_name	salary
10101	Srinivasan	Comp. Sci.	65000
12121	Wu	Finance	90000
15151	Mozart	Music	40000
22222	Einstein	Physics	95000
32343	El Said	History	60000

×

teaches

ID	course_id	sec_id	semester	year
10101	CS-101	1	Fall	2009
10101	CS-315	1	Spring	2010
10101	CS-347	1	Fall	2009
12121	FIN-201	1	Spring	2010
15151	MU-199	1	Spring	2010
22222	PHY-101	1	Fall	2009

Inst.ID	name	dept_name	salary	teaches.ID	course_id	sec_id	semester	year
10101	Srinivasan	Comp. Sci	65000	10101	CS-101	1	Fall	2009
10101	Srinivasan	Comp. Sci	65000	10101	CS-315	1	Spring	2010
10101	Srinivasan	Comp. Sci	65000	10101	CS-347	1	Fall	2009
10101	Srinivasan	Comp. Sci	65000	12121	FIN-201	1	Spring	2010
10101	Srinivasan	Comp. Sci	65000	15151	MU-199	1	Spring	2010
10101	Srinivasan	Comp. Sci	65000	22222	PHY-101	1	Fall	2009
...	...	...	...	...	...	...	...	...
...	...	...	...	...	...	...	...	...
12121	Wu	Finance	90000	10101	CS-101	1	Fall	2009
12121	Wu	Finance	90000	10101	CS-315	1	Spring	2010
12121	Wu	Finance	90000	10101	CS-347	1	Fall	2009
12121	Wu	Finance	90000	12121	FIN-201	1	Spring	2010
12121	Wu	Finance	90000	15151	MU-199	1	Spring	2010
12121	Wu	Finance	90000	22222	PHY-101	1	Fall	2009
...	...	...	...	...	...	...	...	...
...	...	...	...	...	...	...	...	...

# Attribute Name Clashes

---

Two tables can have attributes with the same name.

Beers (**name**, manf)

Bars (**name**, addr, license)

Problem: this ambiguity can lead to confusion if you write a query involving two tables with common column names:

if two same names appear in the WHERE clause

```
SELECT Bars.name  
FROM Bars, Beers  
WHERE name = name;
```

SQL: “**ERROR: Ambiguous name column**”.

# Qualified Column Names

---

Solution: **disambiguate** attributes by specifying the relation name (giving a **qualified name** of a column)

- e.g., Bars.name means the column name from table Bars

We typically qualify a column name to specify the table which the column comes from. (see previous example below)

```
SELECT Bars.name  
FROM Bars, Beers  
WHERE Bars.name = Beers.name;
```



# Qualified Column Names

---

Question: can I use **qualified names** even if there is no ambiguity?

```
SELECT Sells.beer  
FROM Sells  
WHERE Sells.price > 3.00;
```

# Table Name Clashes

---

The **relation-dot-attribute** convention doesn't help if we use the same relation twice in SELECT.

To handle this, we need to define new names for each “instance” of the relation in the FROM clause.

**Example: Find pairs of beers by the same manufacturer.**

Note: we should avoid:

- pairing a beer with itself e.g., (New, New)
- same pairs with different order e.g., (New, Old) (Old, New)

# Table Name Clashes

---

## Beers:

Name	Manf
80/-	Caledonian
Bigfoot Barley Wine	Sierra Nevada
Burraborang Bock	George IV Inn
Crown Lager	Carlton
Fosters Lager	Carlton
Invalid Stout	Carlton
Melbourne Bitter	Carlton
New	Toohey's
Old	Toohey's
Old Admiral	Lord Nelson
Pale Ale	Sierra Nevada
Premium Lager	Cascade
Red	Toohey's
Sheaf Stout	Toohey's
Sparkling Ale	Cooper's
Stout	Cooper's
Three Sheets	Lord Nelson
Victoria Bitter	Carlton

# Table Name Clashes

```
SELECT b1.name, b2.name
FROM Beers b1, Beers b2
WHERE b1.manf = b2.manf AND b1.name <
b2.name;
```

NAME	NAME
-----	-----
Crown Lager	Fosters Lager
Crown Lager	Invalid Stout
Fosters Lager	Invalid Stout
Fosters Lager	Melbourne Bitter
....	

## Beers:

Name	Manf
80/-	Caledonian
Bigfoot Barley Wine	Sierra Nevada
Burraborang Bock	George IV Inn
Crown Lager	Carlton
Fosters Lager	Carlton
Invalid Stout	Carlton
Melbourne Bitter	Carlton
New	Toohey's
Old	Toohey's
Old Admiral	Lord Nelson
Pale Ale	Sierra Nevada
Premium Lager	Cascade
Red	Toohey's
Sheaf Stout	Toohey's
Sparkling Ale	Cooper's
Stout	Cooper's
Three Sheets	Lord Nelson
Victoria Bitter	Carlton

# Joins (1)

---

For all instructors who have taught courses, find their names and the course ID of the courses they taught.

instructor

<i>ID</i>	<i>name</i>	<i>dept_name</i>	<i>salary</i>
10101	Srinivasan	Comp. Sci.	65000
12121	Wu	Finance	90000
15151	Mozart	Music	40000
22222	Einstein	Physics	95000
32343	El Said	History	60000
32456	Chen	Physics	35000

teaches

<i>ID</i>	<i>course_id</i>	<i>sec_id</i>	<i>semester</i>	<i>year</i>
10101	CS-101	1	Fall	2009
10101	CS-315	1	Spring	2010
10101	CS-347	1	Fall	2009
12121	FIN-201	1	Spring	2010
15151	MU-199	1	Spring	2010
22222	PHY-101	1	Fall	2009

# Joins (1)

---

For all instructors who have taught courses, find their names and the course ID of the courses they taught.

```
SELECT name, course_id
FROM instructor, teaches
WHERE instructor.ID = teaches.ID
```

instructor

ID	name	dept_name	salary
10101	Srinivasan	Comp. Sci.	65000
12121	Wu	Finance	90000
15151	Mozart	Music	40000
22222	Einstein	Physics	95000
32343	El Said	History	60000
33456	Chen	Physics	35000

teaches

ID	course_id	sec_id	semester	year
10101	CS-101	1	Fall	2009
10101	CS-315	1	Spring	2010
10101	CS-347	1	Fall	2009
12121	FIN-201	1	Spring	2010
15151	MU-199	1	Spring	2010
22222	PHY-101	1	Fall	2009

# Joins (2)

---

Find instructor names and the courses they taught in 2010.

instructor

<i>ID</i>	<i>name</i>	<i>dept_name</i>	<i>salary</i>
10101	Srinivasan	Comp. Sci.	65000
12121	Wu	Finance	90000
15151	Mozart	Music	40000
22222	Einstein	Physics	95000
32343	El Said	History	60000
32456	G. Li	Physics	87000

teaches

<i>ID</i>	<i>course_id</i>	<i>sec_id</i>	<i>semester</i>	<i>year</i>
10101	CS-101	1	Fall	2009
10101	CS-315	1	Spring	2010
10101	CS-347	1	Fall	2009
12121	FIN-201	1	Spring	2010
15151	MU-199	1	Spring	2010
22222	PHY-101	1	Fall	2009

# Joins (2)

Find instructor names and the courses they taught in 2010.

**SELECT** name, course\_id

**FROM** instructor, teaches

**WHERE** instructor.ID = teaches.ID **AND** year = 2010

instructor

ID	name	dept_name	salary
10101	Srinivasan	Comp. Sci.	65000
12121	Wu	Finance	90000
15151	Mozart	Music	40000
22222	Einstein	Physics	95000
32343	El Said	History	60000
32456	G. Smith	Business	87000

teaches

ID	course_id	sec_id	semester	year
10101	CS-101	1	Fall	2009
10101	CS-315	1	Spring	2010
10101	CS-347	1	Fall	2009
12121	FIN-201	1	Spring	2010
15151	MU-199	1	Spring	2010
22222	PHY-101	1	Fall	2009



# Natural Join

Natural join matches tuples with the same values for all common attributes, and retains only one copy of each common column

**SELECT \* FROM instructor NATURAL JOIN teaches;**

instructor

<i>ID</i>	<i>name</i>	<i>dept_name</i>	<i>salary</i>
10101	Srinivasan	Comp. Sci.	65000
12121	Wu	Finance	90000
15151	Mozart	Music	40000
22222	Einstein	Physics	95000
32343	El Said	History	60000
45565	Katz	Comp. Sci.	75000
76766	Crick	Biology	72000

teaches

<i>ID</i>	<i>course_id</i>	<i>sec_id</i>	<i>semester</i>	<i>year</i>
10101	CS-101	1	Fall	2009
10101	CS-315	1	Spring	2010
10101	CS-347	1	Fall	2009
12121	FIN-201	1	Spring	2010
15151	MU-199	1	Spring	2010
22222	PHY-101	1	Fall	2009

<i>ID</i>	<i>name</i>	<i>dept_name</i>	<i>salary</i>	<i>course_id</i>	<i>sec_id</i>	<i>semester</i>	<i>year</i>
10101	Srinivasan	Comp. Sci.	65000	CS-101	1	Fall	2009
10101	Srinivasan	Comp. Sci.	65000	CS-315	1	Spring	2010
10101	Srinivasan	Comp. Sci.	65000	CS-347	1	Fall	2009
12121	Wu	Finance	90000	FIN-201	1	Spring	2010
15151	Mozart	Music	40000	MU-199	1	Spring	2010
22222	Einstein	Physics	95000	PHY-101	1	Fall	2009
32343	El Said	History	60000	HIS-351	1	Spring	2010
45565	Katz	Comp. Sci.	75000	CS-101	1	Spring	2010
45565	Katz	Comp. Sci.	75000	CS-319	1	Spring	2010
76766	Crick	Biology	72000	BIO-101	1	Summer	2009
76766	Crick	Biology	72000	BIO-301	1	Summer	2010

# Danger of Natural Join

---

List the names of instructors along with the titles of courses that they teach. This is an **incorrect version**:

```
SELECT name, title
FROM instructor
NATURAL JOIN teaches
NATURAL JOIN course;
```

*instructor*

<b>instructor_id</b>	<b>name</b>	<b>dept_name</b>	<b>salary</b>
8	ABC	SEEM	100
7	XYZ	SEEM	120

*teaches*

<b>instruct_id</b>	<b>course_id</b>	<b>s_id</b>	<b>semester</b>	<b>year</b>
7	3550	1	1	2018
8	2100	1	2	2018

*course*

<b>course_id</b>	<b>title</b>	<b>dept_name</b>	<b>credits</b>
3550	DB	SEEM	3
2100	Algo	CSE	3

# Danger of Natural Join

---

- Course.dept\_name and instructor.dept\_name are not related
- Therefore, cannot be assumed to be the same.

*instructor*

<b>instructor_id</b>	<b>name</b>	<b>dept_name</b>	<b>salary</b>
8	ABC	SEEM	100
7	XYZ	SEEM	120

*teaches*

<b>instruct_id</b>	<b>course_id</b>	<b>s_id</b>	<b>semester</b>	<b>year</b>
7	3550	1	1	2018
8	2100	1	2	2018

*course*

<b>course_id</b>	<b>title</b>	<b>dept_name</b>	<b>credits</b>
3550	DB	SEEM	3
2100	Algo	CSE	3

# Correct Natural Join

List the names of instructors  
along with the titles of courses  
that they teach. This is a **correct**  
**version**:

```
SELECT name, title  
FROM instructor, teaches,  
course  
WHERE instructor.ID =  
teaches.ID AND  
teaches.course_id =  
course.course_id;
```

*instructor*

<b>instruct_id</b>	<b>name</b>	<b>dept_name</b>	<b>salary</b>
8	ABC	SEEM	100
7	XYZ	SEEM	120

*teaches*

<b>instruct_id</b>	<b>course_id</b>	<b>s_id</b>	<b>semester</b>	<b>year</b>
7	3550	1	1	2018
8	2100	1	2	2018

*course*

<b>course_id</b>	<b>title</b>	<b>dept_name</b>	<b>credits</b>
3550	DB	SEEM	3
2100	Algo	CSE	3

# JOIN ON

---

books(id, title, translator\_id)

translators(id, name)

```
SELECT books.id, books.title, translators.name AS translator  
FROM books  
JOIN translators ON books.translator_id = translators.id;
```